


AI Foundations

Kun Yuan

Center for Machine Learning Research @ Peking University

Multi-classification


- Multi-classification is very common in real life



[CS231n: Deep Learning for Computer Vision]

Multi-classification

- Face recognition is one of the most successful multi-classification tasks



Elon Musk

Dwayne Johnson

郭达

Adam Lambert

Jason Statham




Predict the score

- We first collect the dataset $\{h_i, y_i\}_{i=1}^N$ where $h_i \in \mathbb{R}^d$ is the feature and $y_i \in \{1, \dots, C\}$ is the label
- We consider a linear model to predict the score of each class

$$s = Xh \in \mathbb{R}^C$$

where $X \in \mathbb{R}^{C \times d}$ is the model parameters to learn and s_i is the score that h belongs to class i




[CS231n: Deep Learning for Computer Vision]

Predict the probability

- Given the score vector s , the probability of each class with the softmax function is as follows


$$p_i = \frac{\exp(s_i)}{\sum_{j=1}^C \exp(s_j)} \in (0, 1)$$



- Given the model parameters X , we can predict the class that feature h belongs to
- But a bad model X will result in incorrect predictions

Cross-entropy

- How to achieve a good model X that can provide accurate predictions?
- We need to train the model so that the discrepancy between ground-truth and predictions are minimized



- How to measure the discrepancy?

- Cross entropy can measure the difference between two distributions $p \in R^d$ and $q \in R^d$

$$H(\mathbf{p}, \mathbf{q}) = - \sum_{j=1}^d p_j \log(q_j)$$

Smaller cross entropy indicates smaller difference between p and q

- Examples:

$$\mathbf{p} = (1, 0, 0, 0) \quad \mathbf{q} = (0.25, 0.25, 0.25, 0.25) \quad \longrightarrow \quad H(\mathbf{p}, \mathbf{q}) = 2$$

$$\mathbf{p} = (1, 0, 0, 0) \quad \mathbf{q} = (0.91, 0.03, 0.03, 0.03) \quad \longrightarrow \quad H(\mathbf{p}, \mathbf{q}) = 0.136$$

Multi-classification: Loss function

- A good model X will minimize the discrepancy between predictions and ground-truth


$$X^* = \arg \min_X \left\{ \frac{1}{N} \sum_{i=1}^N H(\underset{\text{prediction}}{p_i}, \underset{\text{ground-truth}}{q_i}) \right\}, \quad \text{where } p_i = \text{softmax}(\underset{\text{feature of the } i\text{-th sample}}{X h_i})$$

- Once a good model X^* is achieved, we can make predictions as follows


$$p = \text{softmax}(X^* h) \quad \longrightarrow \quad h \text{ belongs to class } j^* = \arg \max_j \{p_j\}$$

Deep neural network (DNN)

- A deep neural network (DNN) typically includes a **feature extractor** and a **classifier**
- Well-trained DNN can make precise predictions



Example: convolutional neural network




Feature extractor

Multi-Classifier

DNN model


- We model DNN as $h(x; \xi) : \mathbb{R}^d \rightarrow \mathbb{R}^c$ that maps input data ξ to a probability $\hat{y}$
 - $x \in \mathbb{R}^d$ is the DNN model parameter to be trained
 - ξ is a random input data sample
 - c is the number of classes



Training DNN can be formulated into an optimization problem

- Define $L(\hat{y}, y) = - \sum_{j=1}^d y_{[j]} \log(\hat{y}_{[j]})$ as the loss function to measure the difference between predictions and the ground-truth label, where $y_{[j]}$ is the j-th element in y
- The model parameter x^* can be achieved by solving the following optimization problem

$$x^* = \arg \min_{x \in \mathbb{R}^d} \left\{ \mathbb{E}_{(\xi, y) \sim \mathcal{D}} \left[L(h(x; \xi), y) \right] \right\}$$


data distribution prediction real label

Training DNN can be formulated into an optimization problem

- The model parameter x^* can be achieved by solving the following optimization problem

$$x^* = \arg \min_{x \in \mathbb{R}^d} \left\{ \mathbb{E}_{(\xi, y) \sim \mathcal{D}} \left[L(\underbrace{h(x; \xi)}_{\text{prediction}}, \underbrace{y}_{\text{real label}}) \right] \right\}$$

data distribution prediction real label

- If we define $\xi = (\xi, y)$ and $F(x; \xi) = L(h(x; \xi), y)$, the above problem becomes


$$x^* = \arg \min_{x \in \mathbb{R}^d} \{ \mathbb{E}_{\xi \sim \mathcal{D}} [F(x; \xi)] \}$$

- Most optimization researchers use the second formulation as the starting point to develop algorithms


DNN model is notoriously difficult to train

- Highly-nonconvex cost functions; cannot find global minima; trapped into local minimum

VGG-56




VGG-110



[Visualizing the loss landscape of neural nets]

DNN model is notoriously difficult to train

- The model size is large, i.e., $x \in \mathbb{R}^d$ is of extremely high dimensions



[Train and deploy large language models on Amazon SageMaker]

DNN model is notoriously difficult to train

- The size of the dataset is huge

Dataset	Sampling prop.	Epochs	Disk size
CommonCrawl	67.0%	1.10	3.3 TB
C4	15.0%	1.06	783 GB
Github	4.5%	0.64	328 GB
Wikipedia	4.5%	2.45	83 GB
Books	4.5%	2.23	85 GB
ArXiv	2.5%	1.06	92 GB
StackExchange	2.0%	1.03	78 GB

} Crawled data from websites; in both high quality and low quality

} High-quality data

Table 1: **Pre-training data.** Data mixtures used for pre-training, for each subset we list the sampling proportion, number of epochs performed on the subset when training on 1.4T tokens, and disk size. The pre-training runs on 1T tokens have the same sampling proportion.

[LLaMA: Open and Efficient Foundation Language Models]

DNN model is notoriously difficult to train

DNN training = Non-convex training + Huge dimensions + Huge dataset

Efficient and **scalable** optimization approaches are in urgent need